

# Database Systems Evolution

UA.DETI.CBD

José Luis Oliveira / Carlos Costa

# Outline

---

- ❖ Why do we need storage system
- ❖ How they evolved along the time
- ❖ Milestone solutions
- ❖ Current landscape

# Thinking about Data Systems

---

- ❖ Many applications today are **data-intensive**, as opposed to **compute-intensive**.
- ❖ Raw CPU power is rarely a limiting factor for these applications
  - bigger problems are usually the **amount** of data, the **complexity** of data, and the **speed** at which it is changing.



# Data systems typically needs to

---

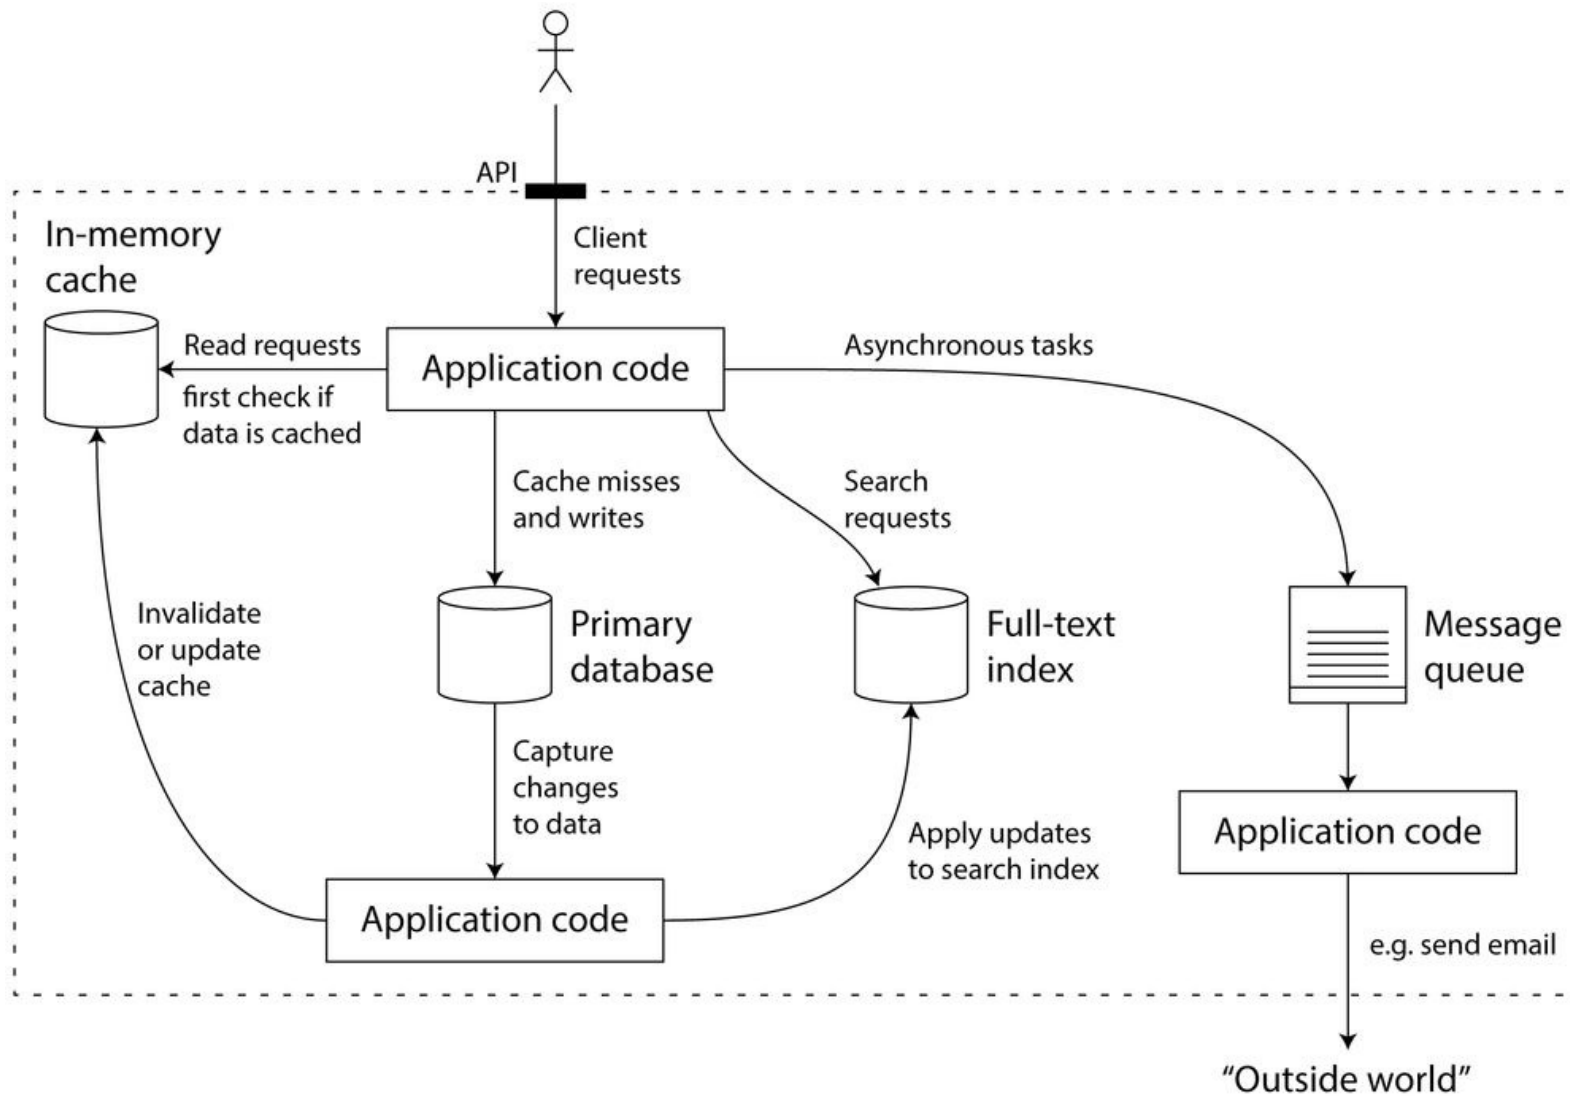
- ❖ Store data so that they, or another application, can find it again later (**databases**).
- ❖ Remember the result of an expensive operation, to speed up reads (**caches**).
- ❖ Allow users to search data by keyword or filter it in various ways (**search indexes**).
- ❖ Send a message to another process, to be handled asynchronously (**message queues**).
- ❖ Observe what is happening, and act on events as they occur (**stream processing**).
- ❖ Periodically crunch a large amount of accumulated data (**batch processing**).

# Thinking about Data Systems

---

- ❖ Increasingly, many applications have wide-ranging requirements
  - Many times, a single tool can no longer meet all of its data processing and storage needs.
- ❖ Instead, the work is broken down into tasks that can be performed efficiently on a single tool,
  - the different tools are stitched together using application code.
- ❖ For example, we may have an application with:
  - a caching layer (e.g. memcached or similar),
  - a full-text search server (e.g. Elasticsearch or Solr),
  - separated from the main database (e.g. MySQL).

# Thinking about Data Systems



# Data Systems – some challenges

---

- ❖ How do you ensure that the data remains correct and complete,
  - even when things go wrong internally?
- ❖ How do you provide consistently good performance to clients,
  - even when parts of your system are degraded?
- ❖ How do you scale to handle an increase in load?
- ❖ What does a good API for the service look like?

# Data Systems – some requirements

---

- ❖ **Reliability:** The system should continue performing the correct function at the desired performance,
  - even in the face of adversity (hardware or software faults, and even human error).
- ❖ **Scalability:** As the system grows (in data volume, traffic volume or complexity), there should be reasonable ways of dealing with that growth.
- ❖ **Maintainability:** Over time, many different people should all be able to work on it productively,
  - Engineering and operations, both maintaining current behavior and adapting the system to new use cases.



# Database Systems

- ❖ A "database" is normally referred as a **set of related data** and its **organization**.
- ❖ A "database management system" (**DBMS**) controls the access to this data.
  - Providing functions that allow writing, searching, updating, retrieving, and removing large quantities of information.



# Brief History of Database Systems

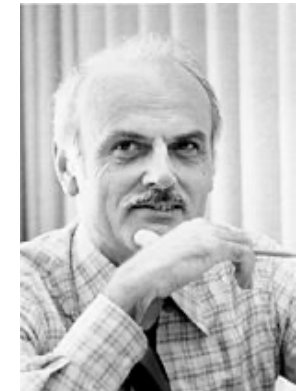
---

## ❖ Pre-relational era (1970's)

- Hierarchical (IMS), Network (Codasyl)
- Many database systems
  - Complex data structures and low-level query language
  - Incompatible, exposing many implementation details

## ❖ **Relational DBMSs (1980s)**

- Edgar F. Codd's relational model in 1970
- Powerful high-level query language
- A few major DB systems dominated the market



## ❖ Object-Oriented DBMSs (1990s)

- Motivated by “mismatch” between RDBMS and OO PL
- Persistent types in C++, Java or Small Talk
- Issues: Lack of high level QL, no standards, performance

# Brief History of Database Systems

---

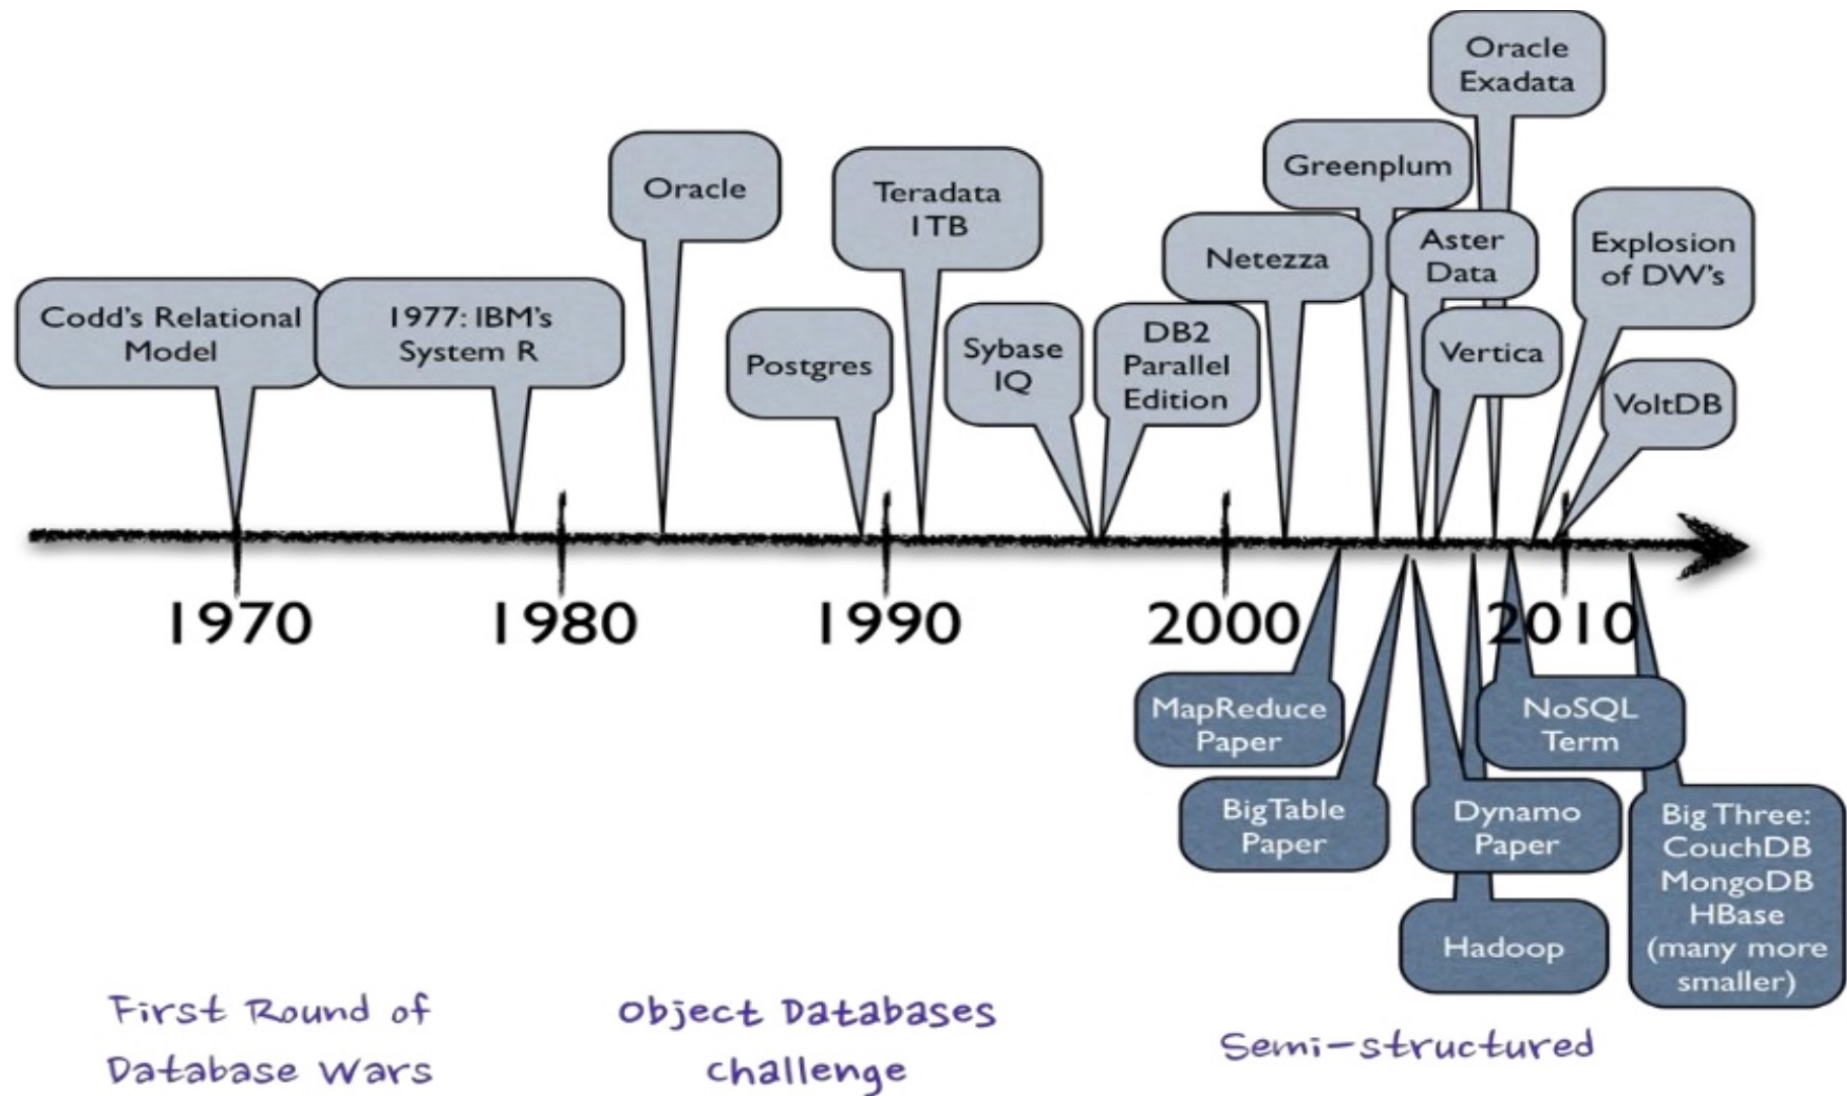
- ❖ Object-relational DBMS (OR-DBMS) (1990s)
  - Relational DBMS vendors' answer to OO
  - User-defined types, functions (spatial, multimedia) Nested tables
  - SQL: 1999 (2003) standards. Plus performance.
- ❖ XML/DBMS (2000s)
  - Web and XML are merging
  - Native support of XML through ORDBMS extension or native XML DBMS
- ❖ Data analytics system (DSS) (2000s)
  - **Data warehousing and OLAP**

# Brief History of Database Systems

---

- ❖ Data stream management systems (2000s)
  - Continuous query against data streams
- ❖ The era of big data (mid 2000-now):
  - **Big data**: datasets that grow so large (terabytes to petabytes) that they become awkward to work with traditional DBMS
  - Parallel DBMSs continue to push the scale of data
  - **MapReduce** dominates on Web data analysis
  - **NoSQL** (not only SQL) is fast growing

# Database Evolution Timeline



# Database Systems Landscape

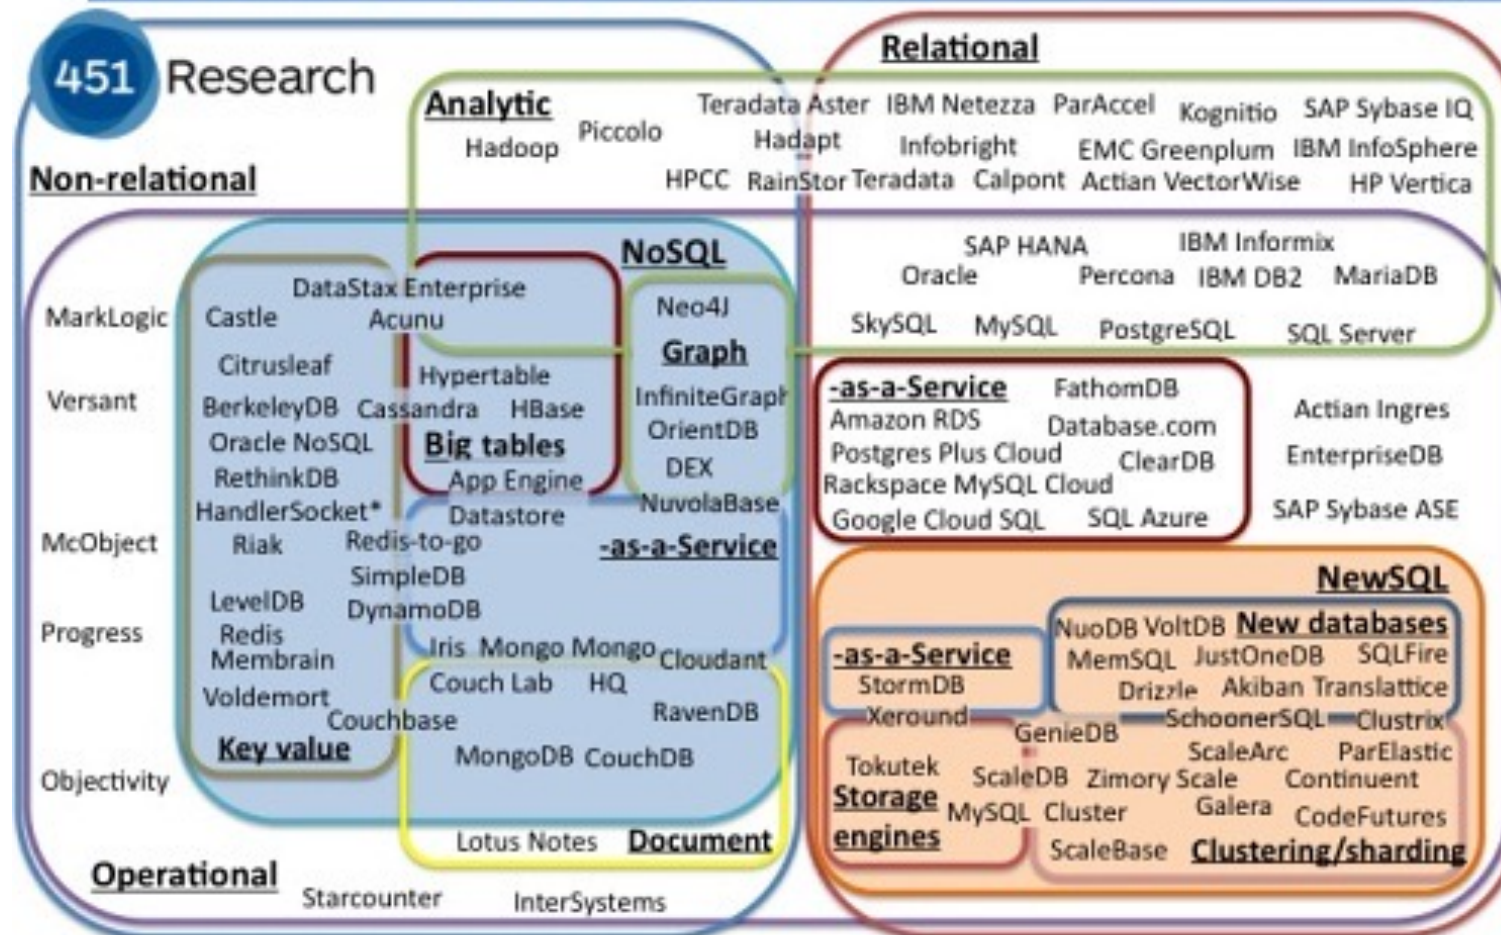
---





# Database Systems Landscape

## The evolving database landscape



© 2012 by The 451 Group. All rights reserved



# Data Platforms Landscape Map – February 2014

451 Research

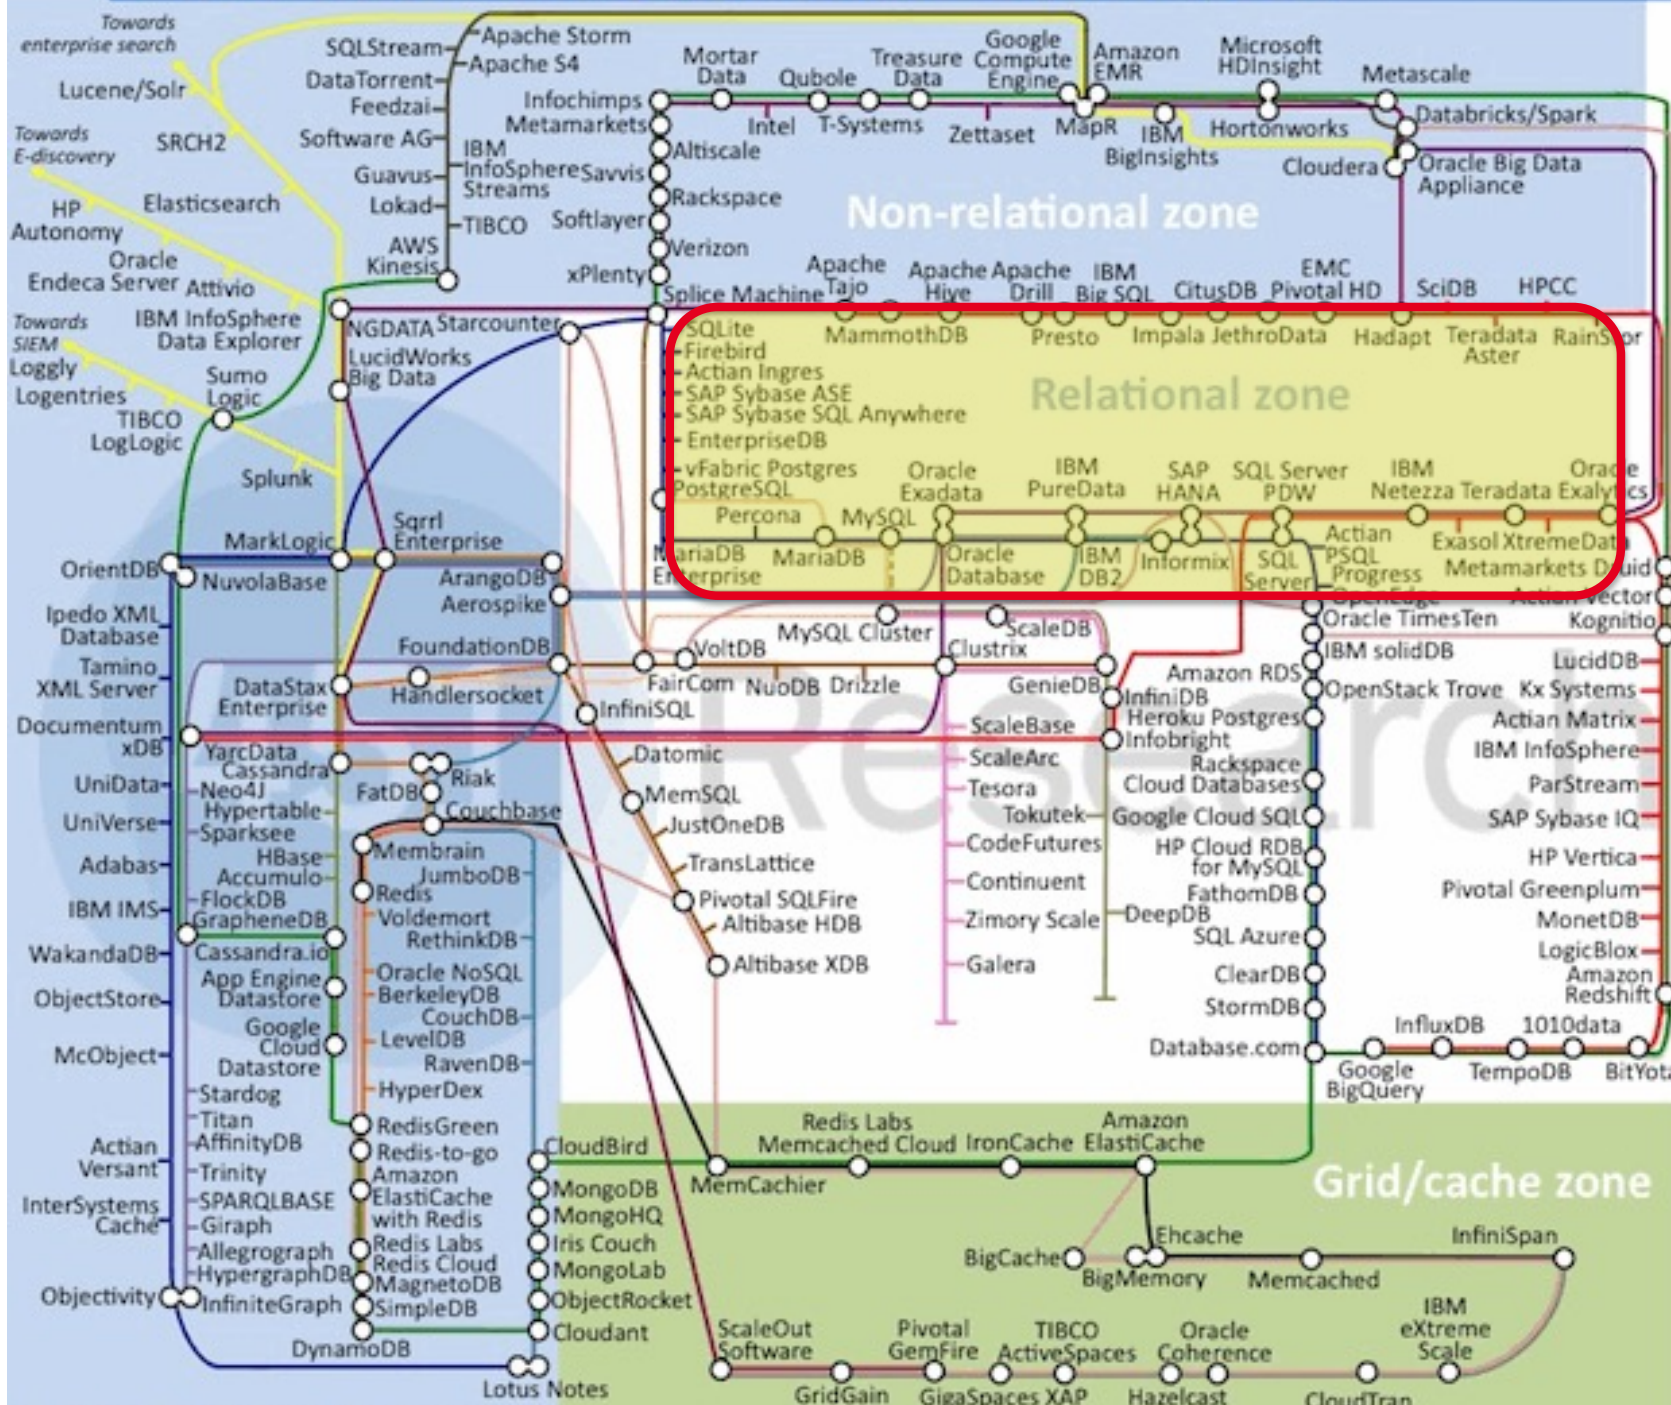
## Key:

- General purpose
- Specialist analytic
- as-a-Service
- NoSQL extension
- BigTables
- Graph
- Document
- Key value stores
- Key value direct access
- Hadoop
- NewSQL extension
- MySQL storage engines
- Advanced clustering/sharding
- New SQL databases
- Data caching
- Data grid
- Search
- Appliances
- Off-heap memory
- In-memory
- Stream processing

Non-relational zone

Relational zone

Grid/cache zone



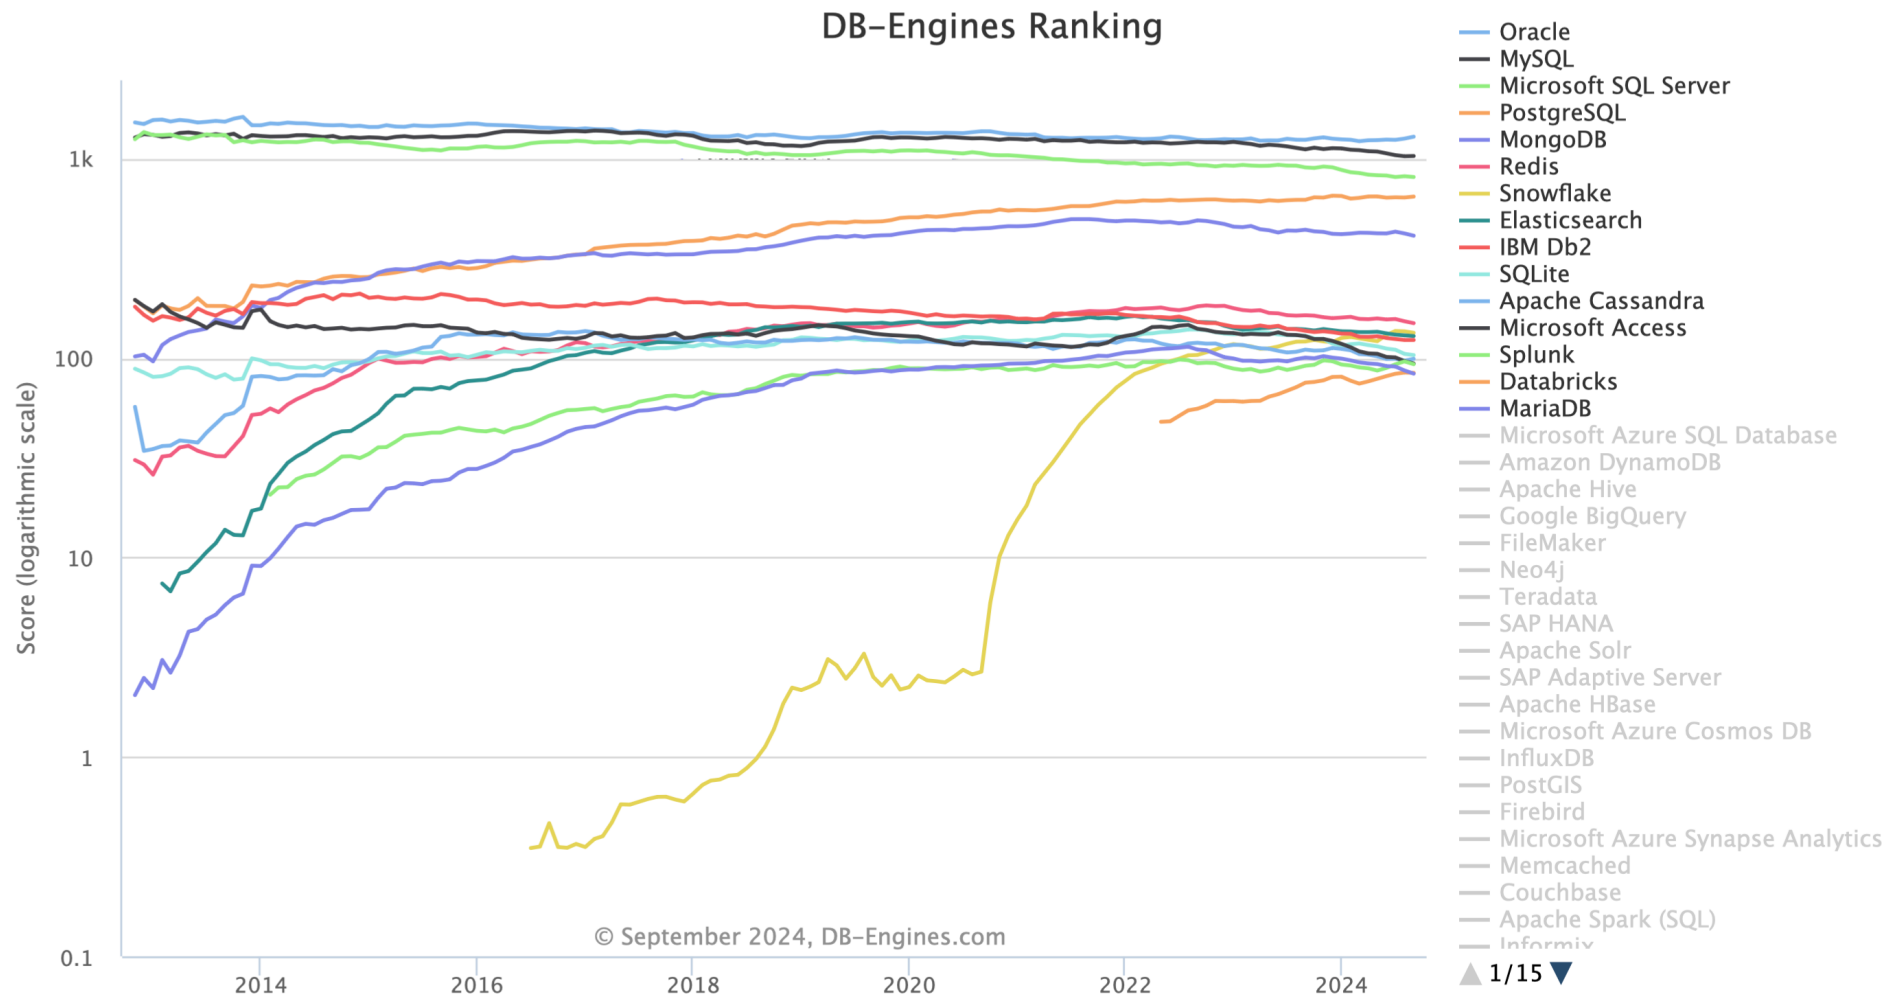


# Database Systems Landscape

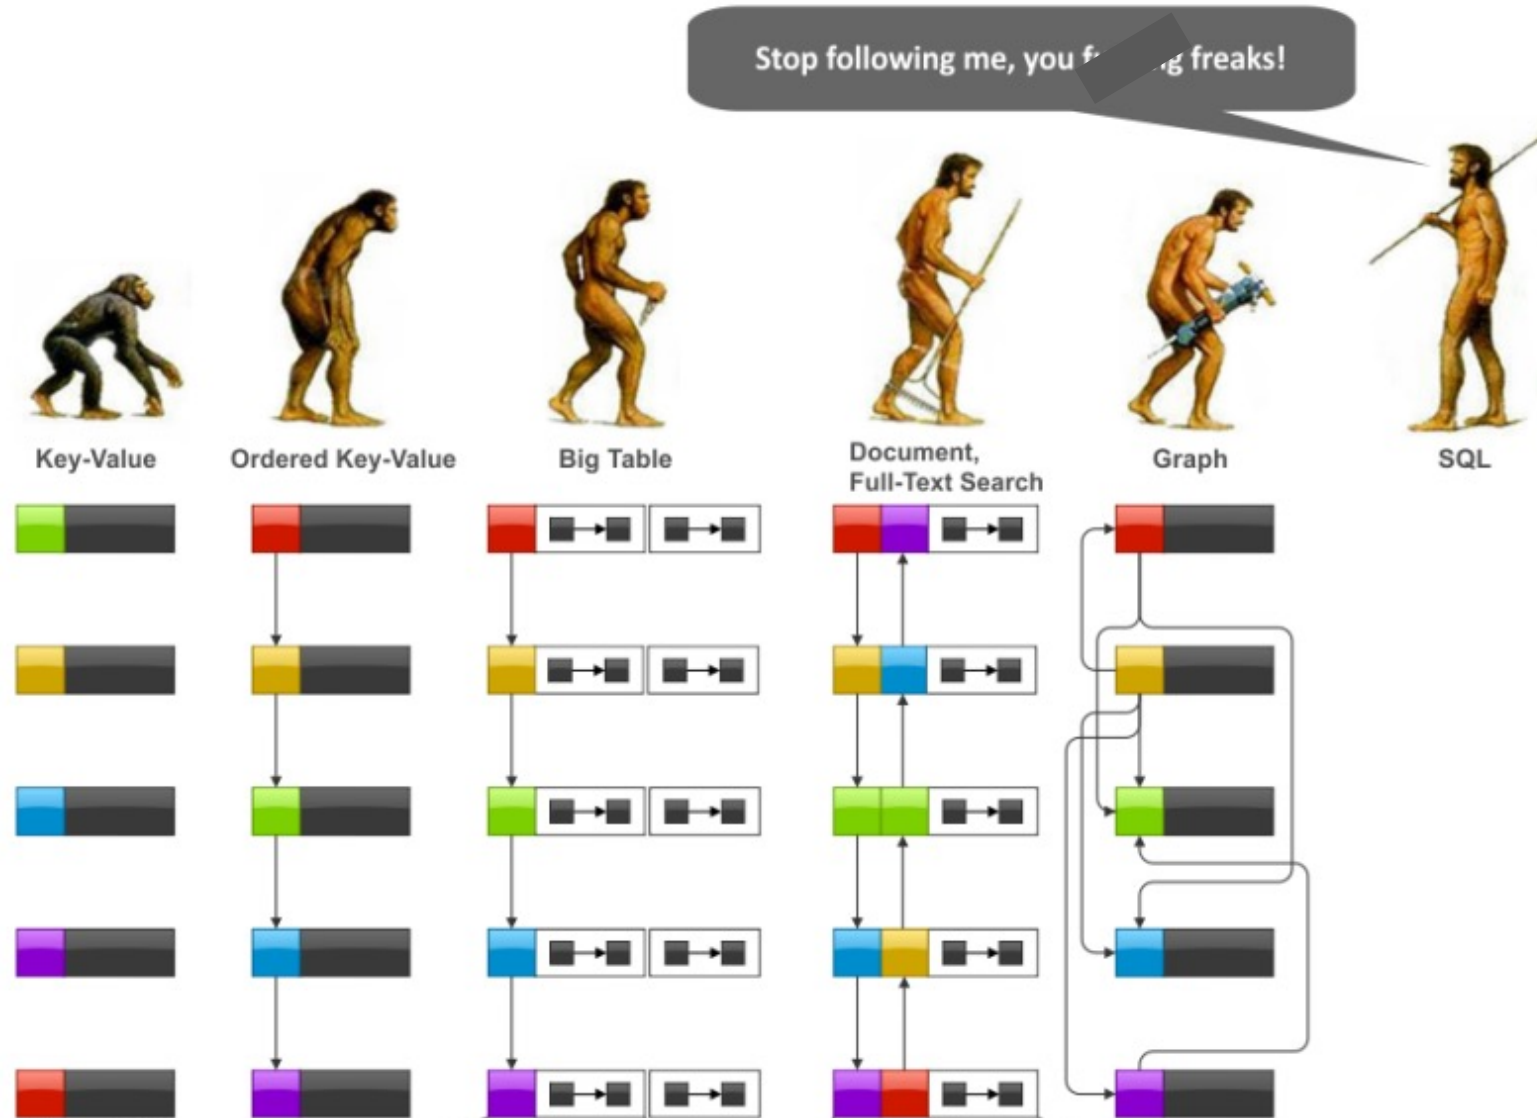
423 systems in ranking, September 2024

Rank			DBMS	Database Model	Score		
Sep 2024	Aug 2024	Sep 2023			Sep 2024	Aug 2024	Sep 2023
1.	1.	1.	Oracle +	Relational, Multi-model ⓘ	1286.59	+28.11	+45.72
2.	2.	2.	MySQL +	Relational, Multi-model ⓘ	1029.49	+2.63	-82.00
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model ⓘ	807.76	-7.41	-94.45
4.	4.	4.	PostgreSQL +	Relational, Multi-model ⓘ	644.36	+6.97	+23.61
5.	5.	5.	MongoDB +	Document, Multi-model ⓘ	410.24	-10.74	-29.18
6.	6.	6.	Redis +	Key-value, Multi-model ⓘ	149.43	-3.28	-14.26
7.	7.	↑ 11.	Snowflake +	Relational	133.72	-2.25	+12.83
8.	8.	↓ 7.	Elasticsearch	Search engine, Multi-model ⓘ	128.79	-1.04	-10.20
9.	9.	↓ 8.	IBM Db2	Relational, Multi-model ⓘ	123.05	+0.04	-13.67
10.	10.	↓ 9.	SQLite +	Relational	103.35	-1.44	-25.85
11.	11.	↑ 12.	Apache Cassandra +	Wide column, Multi-model ⓘ	98.94	+1.94	-11.11
12.	12.	↓ 10.	Microsoft Access	Relational	93.76	-2.61	-34.81
13.	13.	↑ 14.	Splunk	Search engine	93.02	-3.08	+1.63
14.	↑ 15.	↑ 17.	Databricks +	Multi-model ⓘ	84.24	-0.22	+9.06
15.	↓ 14.	↓ 13.	MariaDB +	Relational, Multi-model ⓘ	83.44	-3.09	-17.01
16.	16.	↓ 15.	Microsoft Azure SQL Database	Relational, Multi-model ⓘ	72.95	-2.08	-9.78
17.	17.	↓ 16.	Amazon DynamoDB +	Multi-model ⓘ	70.06	+1.15	-10.85
18.	↑ 19.	18.	Apache Hive	Relational	53.07	-2.17	-18.76
19.	↓ 18.	↑ 20.	Google BigQuery +	Relational	52.67	-2.86	-3.80
20.	20.	↑ 21.	FileMaker	Relational	45.20	-1.47	-8.40
21.	21.	↑ 23.	Neo4j +	Graph	42.68	-1.22	-7.71
22.	↑ 23.	↓ 19.	Teradata	Relational, Multi-model ⓘ	41.47	-0.78	-18.86

# Database Systems Landscape



# Database Systems Landscape



# Resources

---

- ❖ Martin Kleppmann, ***Designing Data-Intensive Applications***, O'Reilly Media, Inc., 2017.
- ❖ Pramod J Sadalage and Martin Fowler, ***NoSQL Distilled*** Addison-Wesley, 2012.
- ❖ Eric Redmond, Jim R. Wilson. ***Seven databases in seven weeks***, Pragmatic Bookshelf, 2012.
- ❖ Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom, ***Database systems: the complete book*** (2nd Ed.), Pearson Education, 2009.