

Distributed Data Processing

UA.DETI.CBD


José Luis Oliveira / Carlos Costa

Introduction

- ❖ Large-Scale Data Processing
 - aim to use 1000s of CPUs, but with simplified managing
- ❖ MapReduce
- ❖ Apache Hadoop
- ❖ To start:
 - <https://www.youtube.com/watch?v=9s-vSeWej1U>



Context

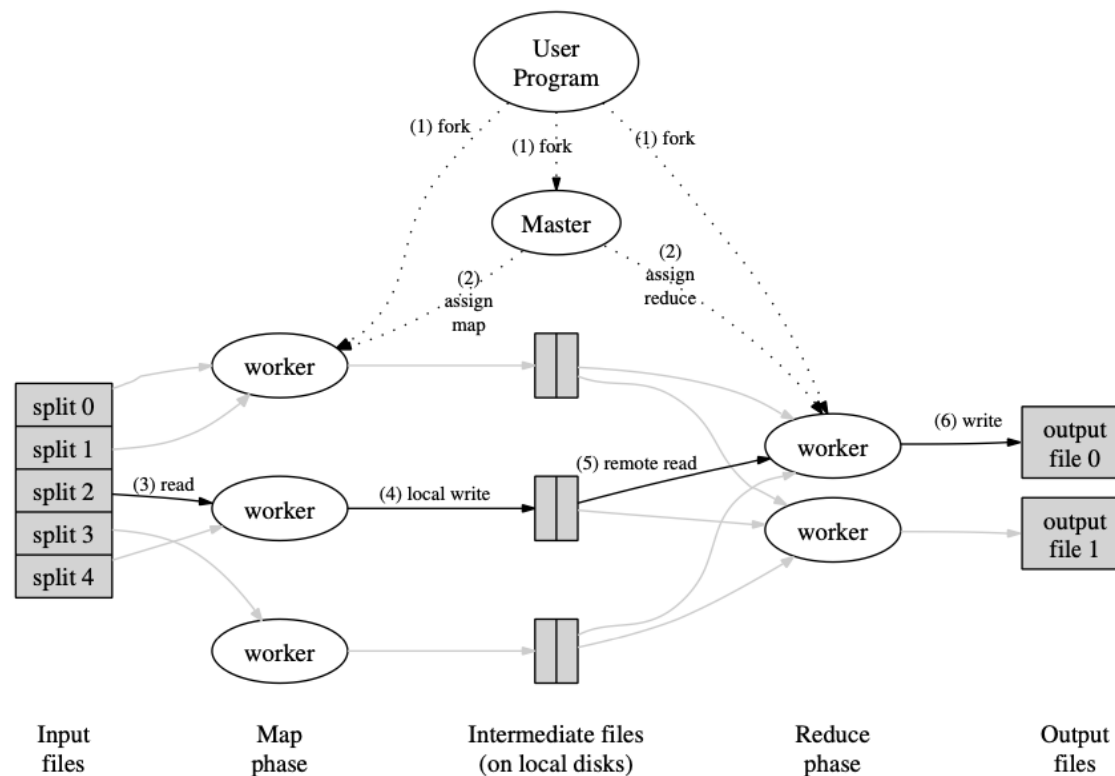


- ❖ Big Data
- ❖ Distributed
 - Not possible to store it in a single place
- ❖ Transmission
 - High Latency

What is MapReduce?

❖ "MapReduce is a programming model and an associated implementation for processing and generating large data sets"

- Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: Simplified data processing on large clusters." (2004). [Google]



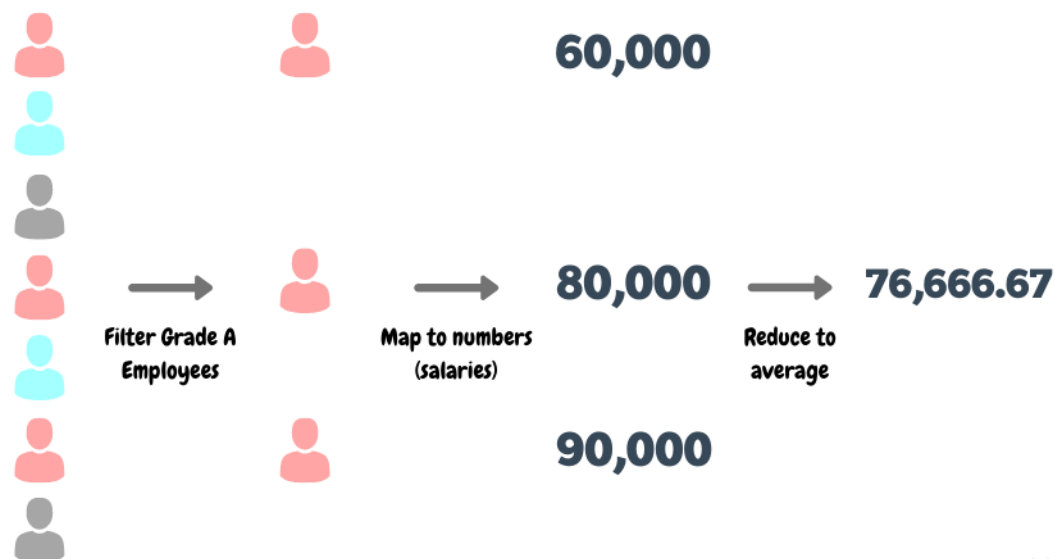
What is MapReduce?

- ❖ Terms are borrowed from Functional Language (e.g., Lisp)
- ❖ Example: Sum of squares
 - **map** square '(1 2 3 4))
Output: (1 4 9 16)
 - *processes each record sequentially and independently*
 - **reduce** + '(1 4 9 16))
(+ 16 (+ 9 (+ 4 1)))
Output: 30
 - *processes set of all records in batches*
- ❖ This concept has been reused in several programming languages and computational tools

Examples: Java stream

```
List<String> values = Arrays.asList("1","2","3","4","5","6","7");  
int soma = values.stream  
    .map(num->Integer.parseInt(num))  
    .reduce(0, Integer::sum);
```

```
OptionalDouble avgSalary = employees.stream()  
    .filter((e) -> e.grade == 'A')           // filter 'A' grade employees  
    .mapToInt((e) -> e.salary)               // get IntStream of salaries  
    .average();                               // get average
```



nextptr

Examples: MongoDB

Database:

```
db.orders.insertMany([
  { _id: 1, cust_id: "Ant O. Knee", ord_date: new Date("2020-03-01"), price: 25,
    items: [ { sku: "oranges", qty: 5, price: 2.5 }, { sku: "apples", qty: 5, price: 2.5 } ],
    status: "A" },
  { _id: 2, cust_id: "Ant O. Knee", ord_date: new Date("2020-03-08"), price: 70, items: [
    { sku: "oranges", qty: 8, price: 2.5 }, { sku: "chocolates", qty: 5, price: 10 } ], status:
    "A" },
  ...
])
```

Return the Total Price Per Customer:

```
var mapFunction1 = function() {
  emit(this.cust_id, this.price);
};
var reduceFunction1 = function(keyCustId, valuesPrices) {
  return Array.sum(valuesPrices);
};
db.orders.mapReduce(
  mapFunction1,
  reduceFunction1,
  { out: "map_reduce_example" }
)
db.map_reduce_example.find().sort( { _id: 1 } )
```

```
{ "_id" : "Ant O. Knee", "value" : 95 }
{ "_id" : "Busby Bee", "value" : 125 }
{ "_id" : "Cam Elot", "value" : 60 }
```

...

MapReduce frameworks

- ❖ Library/Tools that allow easily writing applications that process large amounts of data in parallel
 - distributed across several nodes
- ❖ Good retry/failure semantics
- ❖ Solution Pattern
 - many problems can be modelled in this way
- ❖ Implementations
 - **Hadoop**: the mapper and reducer are each a Java class that implements a particular interface.
 - **Spark**: newer and faster, it processes data in RAM using a concept known as an RDD, Resilient Distributed Dataset.

MapReduce frameworks

❖ Framework...

- Automatic parallelization & distribution
- Fault tolerance
- I/O scheduling
- Monitoring & status updates

❖ Two main tasks:

- Mappers & Reducers

❖ Pipeline

- Splits the input data-set into independent chunks
- Mappers process chunks in a parallel manner
- aggregates the outputs of the maps
- Reducers process the mappers output

MapReduce dataflow

❖ The **map** function


- is called once for every input record to extract (one or more) key-value from the input record

❖ MapReduce framework

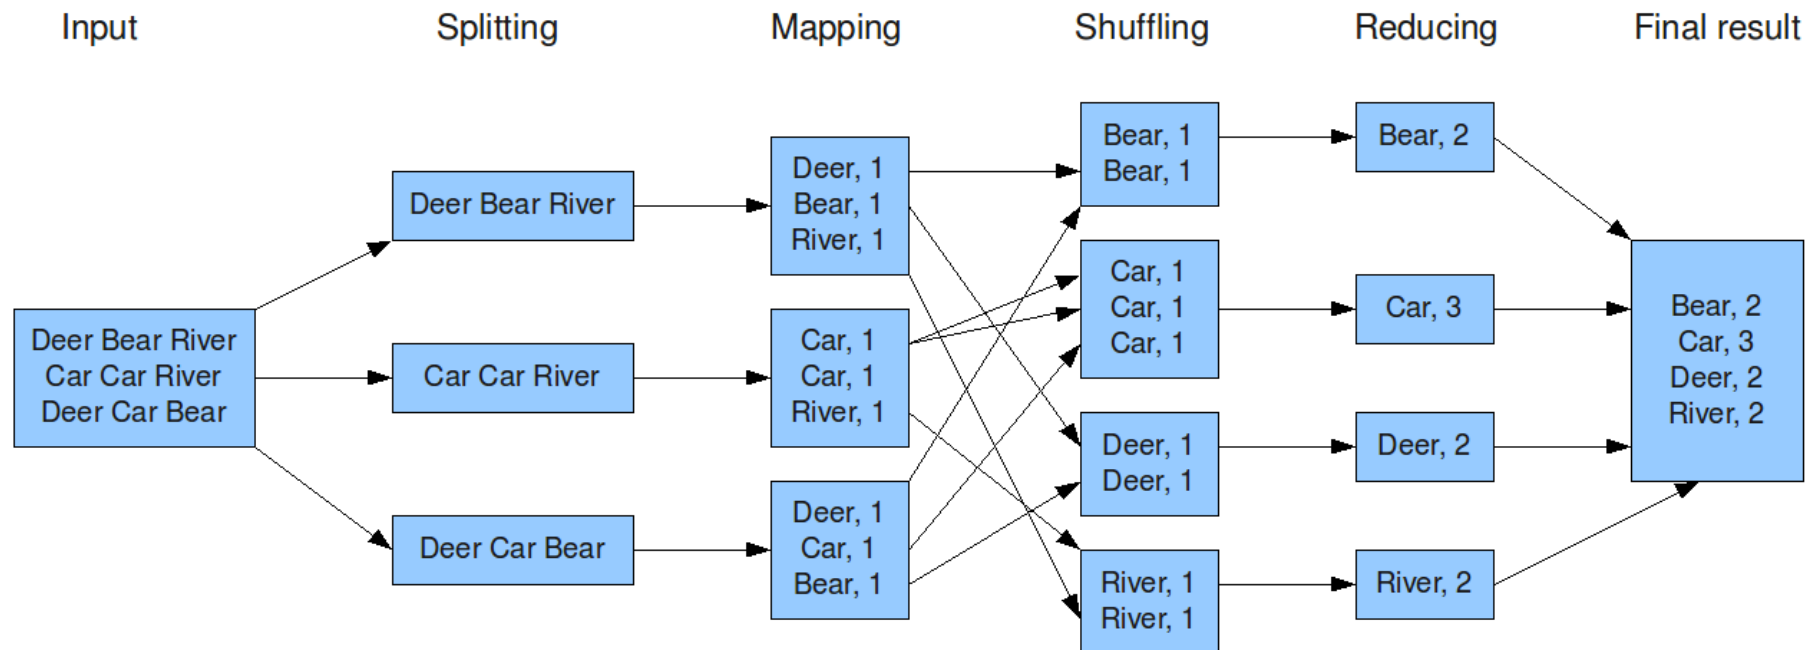
- collects all the key-value pairs with the same key – **shuffle**

❖ The **reduce** function

- iterates over each collection of values with the same key and can produce output records
 - e.g., the number of occurrences of the key



Example: Words count dataflow



Words count – tasks execution

Map Input

Page 1: the weather is good

Page 2: today is good

Page 3: good weather is good




Map output

Worker 1: (the 1), (weather 1), (is 1), (good 1)

Worker 2: (today 1), (is 1), (good 1)

Worker 3: (good 1), (weather 1), (is 1), (good 1)



Reduce output


Worker 1: (the 1)

Worker 2: (is 3)

Worker 3: (weather 2)

Worker 4: (today 1)

Worker 5: (good 4)



Reduce input

Worker 1: (the 1)

Worker 2: (is 1), (is 1), (is 1)

Worker 3: (weather 1), (weather 1)


Worker 4: (today 1)

Worker 5: (good 1), (good 1), (good 1), (good 1)

MapReduce applications

- ❖ words count / histogram
- ❖ distributed search
- ❖ distributed sort
- ❖ web link-graph reversal
- ❖ term-vector per host
- ❖ web access log stats
- ❖ inverted index construction
- ❖ document clustering
- ❖ machine learning
- ❖ statistical machine translation
- ❖ ...

Apache Hadoop



Apache Hadoop Framework

❖ A framework that allows distributed processing of **large data sets across clusters of computers** using simple programming models.

❖ Open-source framework

- Java

- <https://hadoop.apache.org>

❖ **Main components**

- Hadoop Distributed File System (**HDFS**)

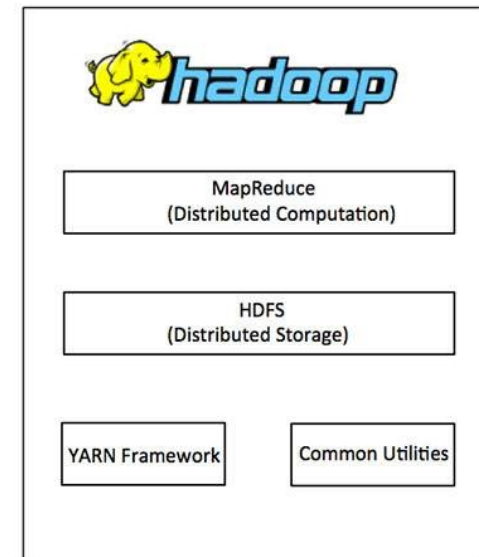
- Distributed, scalable, and portable file system

- Hadoop Yet Another Resource Negotiator (**YARN**)

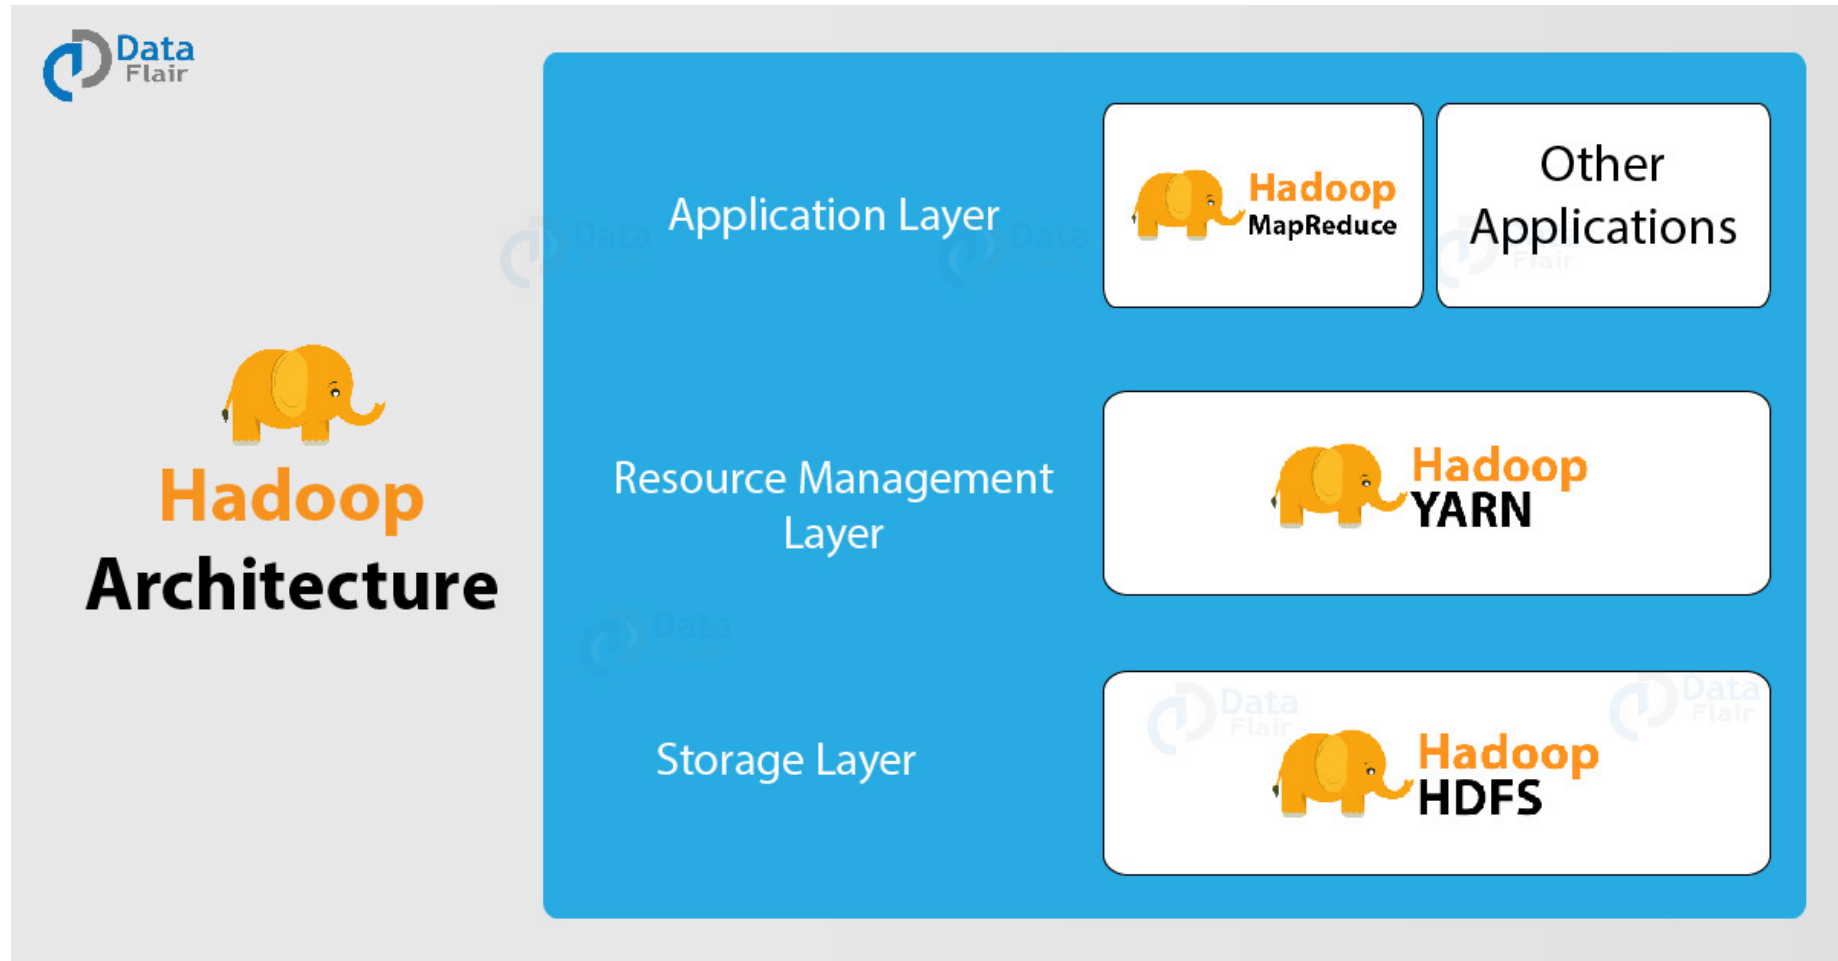
- Hadoop Common Utilities

- Hadoop **MapReduce**

- Implementation of the MapReduce programming model



Apache Hadoop Framework



Hadoop Distributed File System (HDFS)

- ❖ Distributed file system designed to run on commodity hardware.
- ❖ Stores data redundantly on multiple nodes – Fault-tolerant file system
 - 3+ replicas for each block
 - Default Block Size : 128MB
- ❖ Master-Slave architecture
 - NameNode: Master
 - DataNode: Slave
- ❖ Typical usage pattern
 - Huge files (GB to TB)
 - Data is rarely updated
 - Reads and appends are common
 - Usually, random read/write operations are not performed

HDFS – Assumptions and Goals

❖ **Hardware Failure**

- An HDFS instance may consist of +100/+1000 of server machines.

❖ **Streaming Data Access**

- HDFS is designed more for batch processing rather than interactive use. The emphasis is on high throughput of data access rather than low latency of data access.

❖ **Large Data Sets**

- Applications that run on HDFS have large data sets. A typical file in HDFS is gigabytes to terabytes in size.

❖ **Simple Coherency Model**

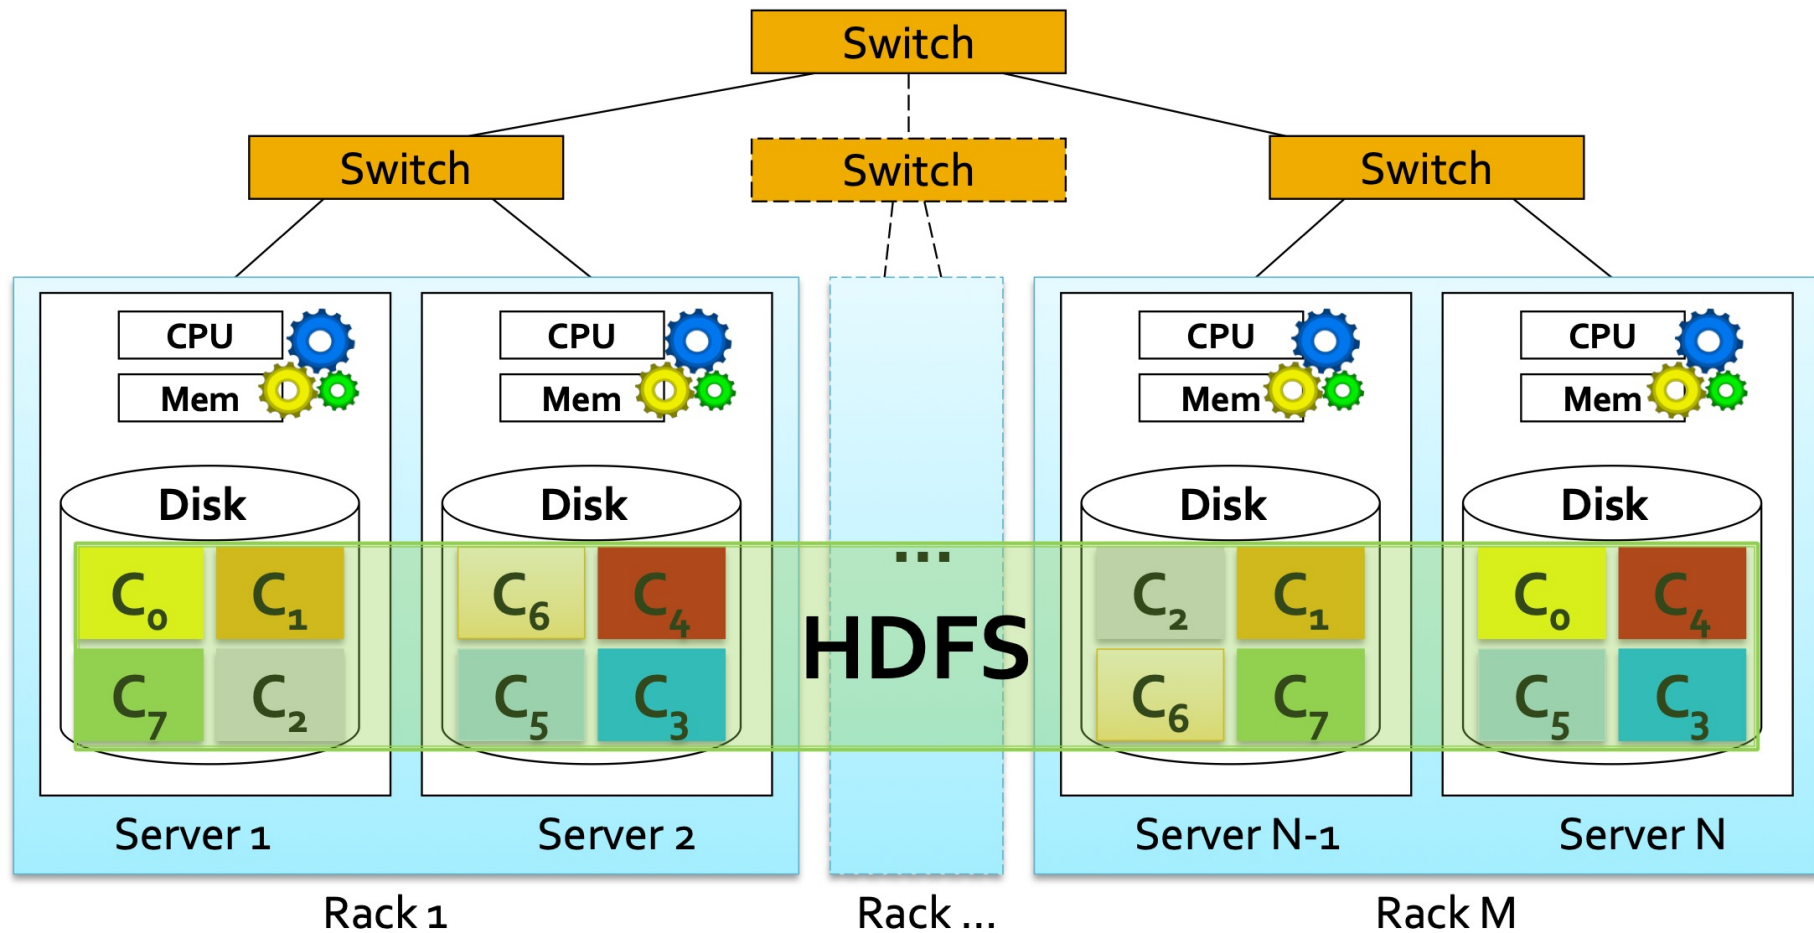
- HDFS applications need a write-once-read-many access model for files. A file once created, written, and closed need not be changed except for appends and truncates.

❖ **“Moving Computation is Cheaper than Moving Data”**

- Especially true when the size of the data set is huge. Minimizes network congestion and increases the throughput of the system.

❖ **Portability Across Heterogeneous Platforms**

HDFS



Example with number of replicas per chunk = 2

HDFS

❖ Master-Slave architecture

- Master: NameNode
- Slave: DataNode

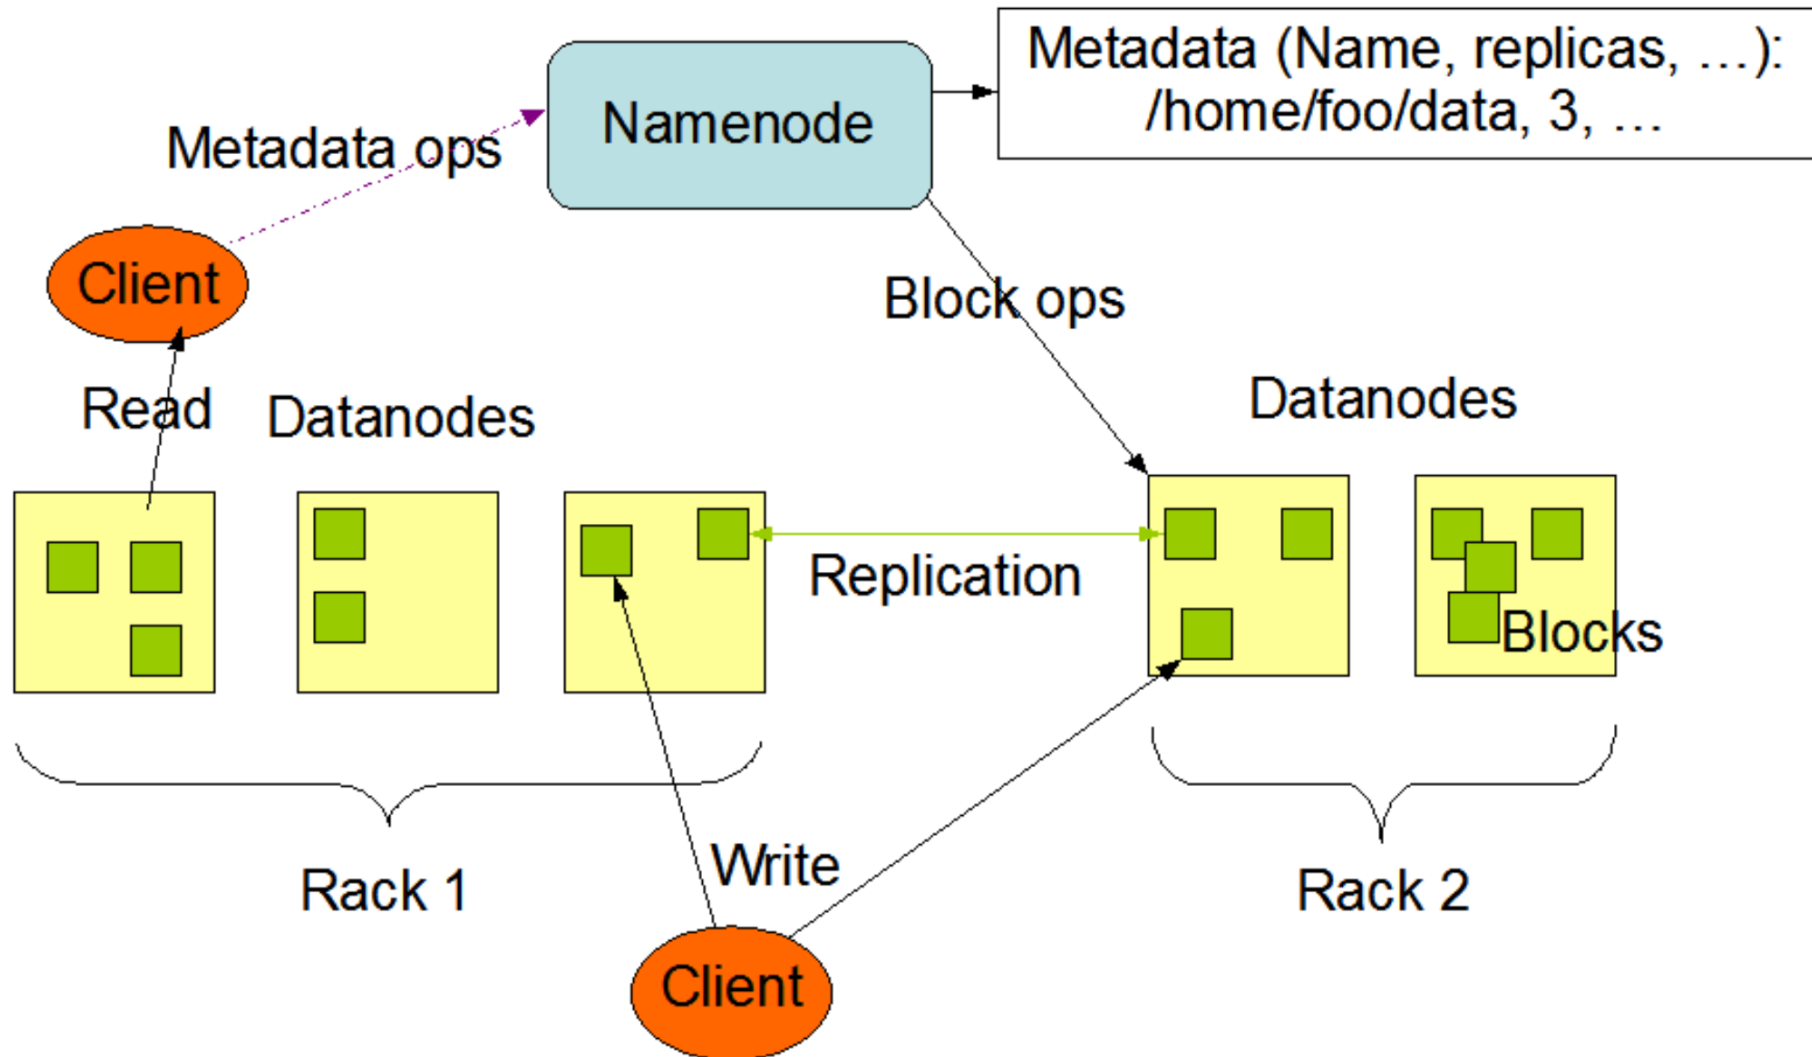
❖ The Master node is a special node/server that

- Stores HDFS metadata
 - e.g., the mapping between the name of a file and the location of its chunks
- Might be replicated

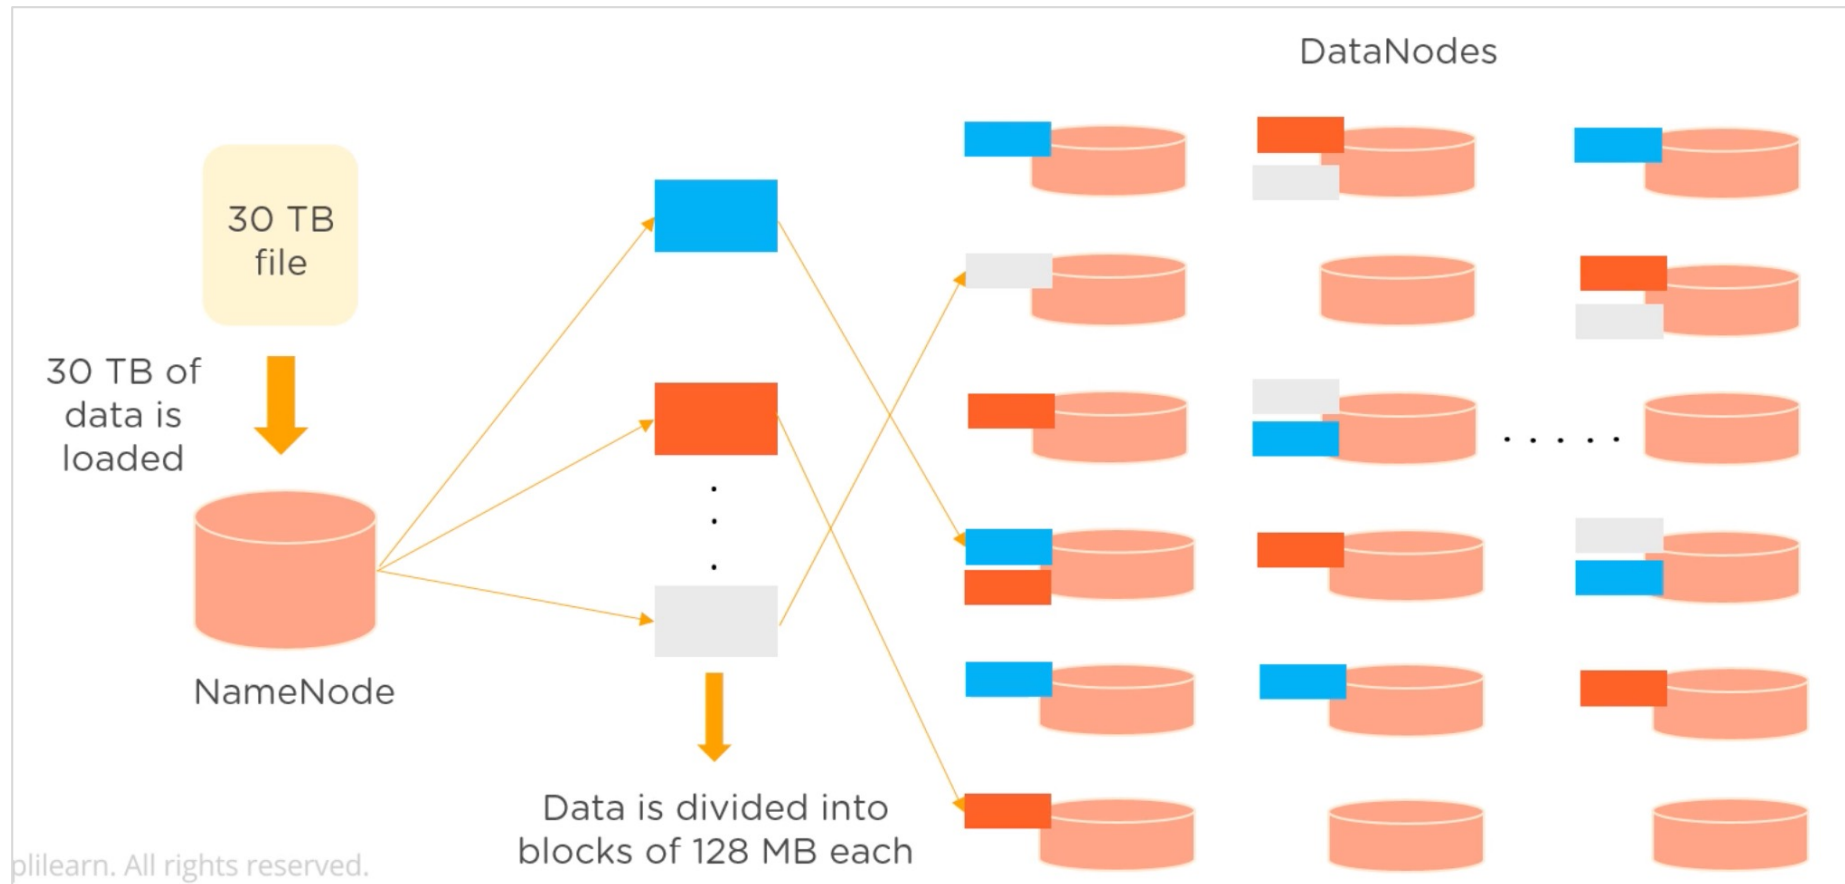
❖ Client applications: file access through HDFS APIs

- Talk to the master node to find data/chuck servers associated with the file of interest
- Connect to the selected chunk servers to access data

Interaction between HDFS components



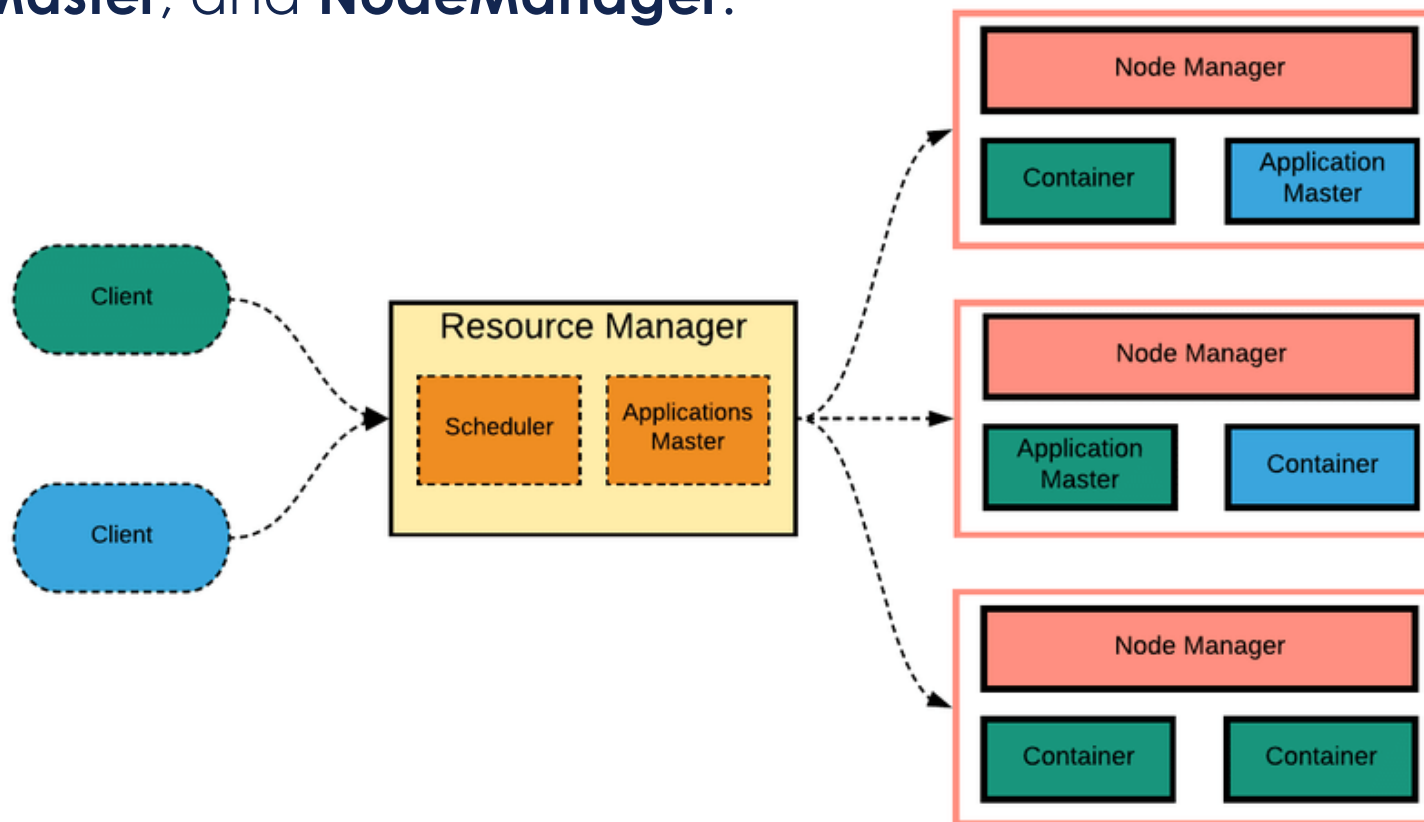
Storage example



<https://www.youtube.com/watch?v=iANBytZ26MI>

Hadoop YARN

- ❖ Resource management and job scheduling.
 - Global **ResourceManager**, per-application **Application Master**, and **NodeManager**.



Hadoop MapReduce

- ❖ MapReduce is the data processing layer of Hadoop.
- ❖ MapReduce job comprises a number of map tasks and reduces tasks.
 - Each task works on a part of data. This distributes the load across the cluster.
 - **Map task** load, parse, transform and filter data.
 - Each **reduce task** works on the sub-set of output from the map tasks (e.g., by applying grouping and aggregation).

Apache Hadoop

- ❖ Install & config

- ❖ Run

 - ~ `sbin ./start-dfs.sh`

- ❖ Basic help for all the commands

 - ~ `hadoop`

- ❖ Distributed file system commands

 - ~ `hadoop fs`

- ❖ Execution of MapReduce jobs

 - ~ `hadoop jar`

HDFS commands

~ **hadoop** fs [generic options]
[-appendToFile <localsrc> ... <dst>]
[-cat [-ignoreCrc] <src> ...]
[-checksum [-v] <src> ...]
[-chgrp [-R] GROUP PATH...]
[-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
[-chown [-R] [OWNER][:[GROUP]] PATH...]
[-copyFromLocal [-f] [-p] [-l] [-d] [-t <thread count>] <localsrc> ... <dst>]
[-copyToLocal [-f] [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
[-count [-q] [-h] [-v] [-t <storage type>]] [-u] [-x] [-e] <path> ...]
[-cp [-f] [-p | -p[topax]] [-d] <src> ... <dst>]
[-createSnapshot <snapshotDir> [<snapshotName>]]
[-deleteSnapshot <snapshotDir> <snapshotName>]
[-df [-h] [<path> ...]]
[-du [-s] [-h] [-v] [-x] <path> ...]
[-expunge [-immediate] [-fs <path>]]
[-find <path> ... <expression> ...]
[-get [-f] [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
...

Hadoop – HDFS Commands

- ❖ Help
 - ~ `hadoop fs -help`
 - ❖ Creating a directory
 - ~ `hadoop fs -mkdir /myhdfsdir`
 - ❖ Listing a directory
 - ~ `hadoop fs -ls`
 - ~ `hadoop fs -ls /myhdfsdir`
 - ❖ Copy a file/dir from local file system to HDFS
 - ^ `hadoop fs -put /myhdfsdir/ /user/xpto/`
 - ❖ Get a file/dir from HDFS to local file system
 - ~ `hadoop fs -get /user/xpto/ /myhdfsdir/`
 - ... and many more
- https://www.tutorialspoint.com/hadoop/hadoop_command_reference.htm

Execute a MapReduce job

❖ Run the job

~ **hadoop jar** WordCount.jar /myhdfsdir/myapp/input1 /myhdfsdir/myapp/output1

❖ Retrieve and explore the job result

~ **hadoop fs -copyToLocal** /myhdfsdir/myapp/output1/part-r-00000 result.txt

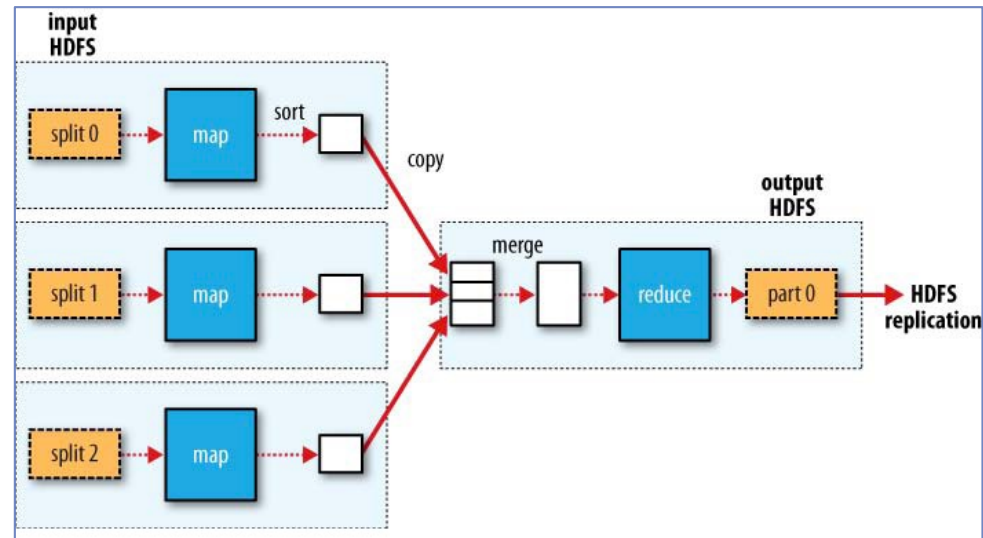
~ **cat** result.txt

❖ Clean the output HDFS directory

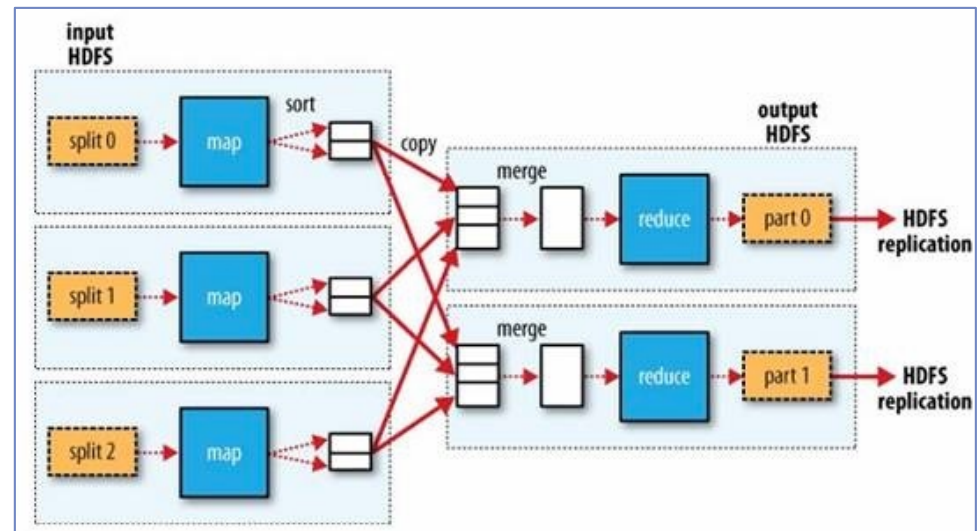
~ **hadoop fs -rm** /myhdfsdir/myapp/output1/

Task execution - distinct configurations

single reduce task



multiple reduce task



Java Interface

❖ Mapper class

- Implementation of the **map function**
- Template parameters
 - KEYIN, VALUEIN** – types of input key-value pairs
 - KEYOUT, VALUEOUT** – types of intermediate key-value pairs
- Intermediate pairs are emitted via **context.write(k, v)**

```
class MyMapper extends Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {  
    @Override  
    public void map(KEYIN key, VALUEIN value, Context context)  
        throws IOException, InterruptedException  
    {  
        // Implementation  
    }  
}
```

Java Interface

❖ Reducer class

- Implementation of the **reduce function**
- Template parameters
 - KEYIN, VALUEIN** – types of intermediate key-value pairs
 - KEYOUT, VALUEOUT** – types of output key-value pairs
- Output pairs are emitted via **context.write(k, v)**

```
class MyReducer extends Reducer<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {  
    @Override  
    public void reduce(KEYIN key, Iterable<VALUEIN> values, Context context)  
        throws IOException, InterruptedException  
    {  
        // Implementation  
    }  
}
```

MapReduce – WordCount example (1)

```
public class WordCount extends Configured implements Tool {
    private static IntWritable ONE = new IntWritable(1);
    @Override
    public int run(String[] args) throws Exception {
        FileSystem fs = FileSystem.get(getConf());
        Job job = Job.getInstance(getConf());
        job.setJarByClass(WordCount.class);
        job.setJobName("WordCount");
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setMapperClass(Map.class); ← Mapper class
        job.setCombinerClass(Reduce.class);
        job.setReducerClass(Reduce.class); ← Reducer class
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.setInputPaths(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        if (fs.exists(new Path(args[1])))
            fs.delete(new Path(args[1]), true);
        boolean success = job.waitForCompletion(true);
        return success ? 0 : 1;
    }
}
```

Configurations

MapReduce – WordCount example (2)

```
...
public static void main(String[] args) throws Exception {
    int ret = ToolRunner.run(new WordCount(), args);
    System.exit(ret);
}

public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException {

        String line = value.toString();
        String[] words = line.split("[\\t\\n.,:; ?!-/()\\[\\]\\\\\\\"' ]+");
        for (String word : words) {
            if (word.trim().length() > 0) {
                Text text = new Text();
                text.set(word);
                context.write(text, ONE);
            }
        }
    }
}
// end class Map
```

Mapper

MapReduce – WordCount example (3)

```
...

public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable>
{
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values)
            sum += val.get();
        context.write(key, new IntWritable(sum));
    }
} // end class Reduce

} // end class WordCount
```

Reducer

MapReduce – Run the example

Compile

```
> $ javac -classpath $CLASSPATH -d WordCountDir WordCount.java
```

Create JAR

```
> $ jar -cvf WordCount.jar -C WordCountDir/ .
```

Run WordCount App

```
> $ hadoop jar WordCount.jar /in/test.txt /out/wc
```

See the Result

```
> $ hadoop fs -cat /out/wc/part-r-00000
```

ABOUT 1

ACCOUNT 2

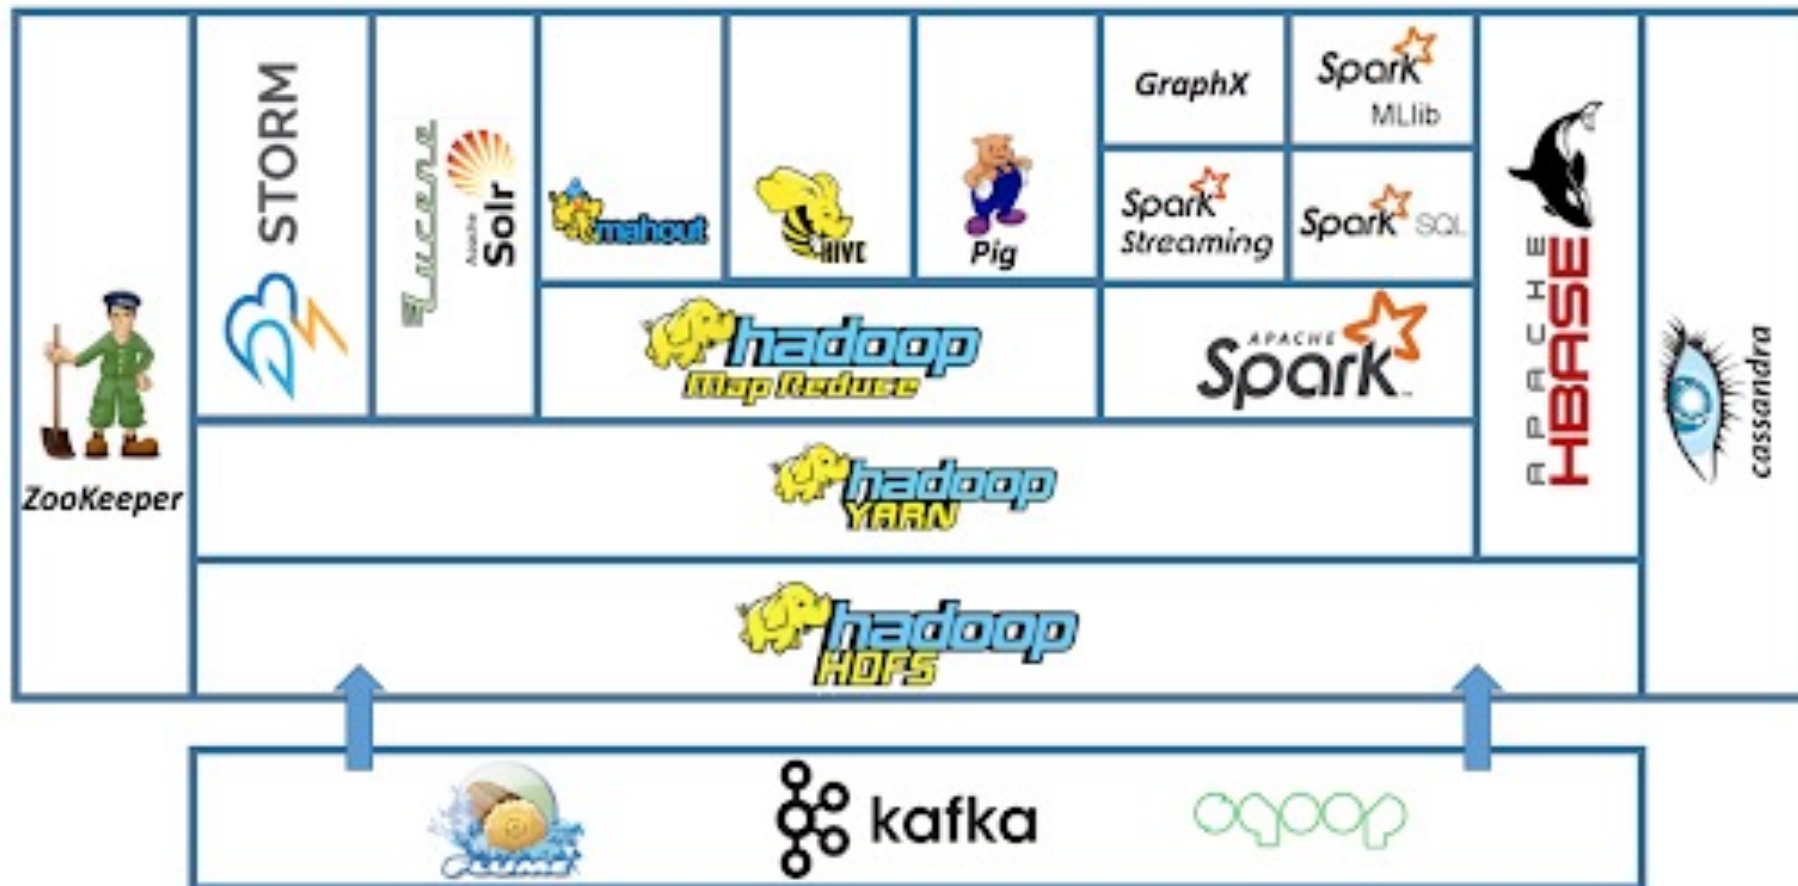
ACTUAL 1

ADDITIONAL 1

ADVANCING 2

...

Hadoop ecosystem



Summary

- ❖ Concept
- ❖ MapReduce framework
 - Programming Model
 - Execution Plan
 - Dataflow
- ❖ MapReduce Applications
- ❖ Apache Hadoop
 - HDFS
 - MapReduce Programming